

# PySpark-analysis-file

November 10, 2020

## 0.1 Migrating from Spark to BigQuery via Dataproc – Part 1

- [Part 1](#): The original Spark code, now running on Dataproc (lift-and-shift).
- [Part 2](#): Replace HDFS by Google Cloud Storage. This enables job-specific-clusters. (cloud-native)
- [Part 3](#): Automate everything, so that we can run in a job-specific cluster. (cloud-optimized)
- [Part 4](#): Load CSV into BigQuery, use BigQuery. (modernize)
- [Part 5](#): Using Cloud Functions, launch analysis every time there is a new file in the bucket. (serverless)

```
[5]: %%writefile spark_analysis.py

import matplotlib
matplotlib.use('agg')

import argparse
parser = argparse.ArgumentParser()
parser.add_argument("--bucket", help="bucket for input and output")
args = parser.parse_args()

BUCKET = args.bucket
```

Overwriting spark\_analysis.py

### 0.1.1 Copy data to HDFS

The Spark code in this notebook is based loosely on the [code](#) accompanying [this post](#) by Dipanjan Sarkar. I am using it to illustrate migrating a Spark analytics workload to BigQuery via Dataproc.

The data itself comes from the 1999 KDD competition. Let's grab 10% of the data to use as an illustration.

### 0.1.2 Reading in data

The data are CSV files. In Spark, these can be read using textFile and splitting rows on commas.

```
[6]: %%writefile -a spark_analysis.py

from pyspark.sql import SparkSession, SQLContext, Row
gcs_bucket="qwiklabs-gcp-02-6236be6b5a47"
```

```

spark = SparkSession.builder.appName("kdd").getOrCreate()
sc = spark.sparkContext
# data_file = "hdfs:///kddcup.data_10_percent.gz"
data_file = "gs://" + gcs_bucket + "/kddcup.data_10_percent.gz"
raw_rdd = sc.textFile(data_file).cache()
raw_rdd.take(5)

```

Appending to spark\_analysis.py

```

[7]: %%writefile -a spark_analysis.py

csv_rdd = raw_rdd.map(lambda row: row.split(","))
parsed_rdd = csv_rdd.map(lambda r: Row(
    duration=int(r[0]),
    protocol_type=r[1],
    service=r[2],
    flag=r[3],
    src_bytes=int(r[4]),
    dst_bytes=int(r[5]),
    wrong_fragment=int(r[7]),
    urgent=int(r[8]),
    hot=int(r[9]),
    num_failed_logins=int(r[10]),
    num_compromised=int(r[12]),
    su_attempted=r[14],
    num_root=int(r[15]),
    num_file_creations=int(r[16]),
    label=r[-1]
))
parsed_rdd.take(5)

```

Appending to spark\_analysis.py

### 0.1.3 Spark analysis

One way to analyze data in Spark is to call methods on a dataframe.

```

[8]: %%writefile -a spark_analysis.py

sqlContext = SQLContext(sc)
df = sqlContext.createDataFrame(parsed_rdd)
connections_by_protocol = df.groupBy('protocol_type').count().orderBy('count',
↪ascending=False)
connections_by_protocol.show()

```

Appending to spark\_analysis.py

Another way is to use Spark SQL

```
[9]: %%writefile -a spark_analysis.py

df.registerTempTable("connections")
attack_stats = sqlContext.sql("""
    SELECT
        protocol_type,
        CASE label
            WHEN 'normal.' THEN 'no attack'
            ELSE 'attack'
        END AS state,
        COUNT(*) as total_freq,
        ROUND(AVG(src_bytes), 2) as mean_src_bytes,
        ROUND(AVG(dst_bytes), 2) as mean_dst_bytes,
        ROUND(AVG(duration), 2) as mean_duration,
        SUM(num_failed_logins) as total_failed_logins,
        SUM(num_compromised) as total_compromised,
        SUM(num_file_creations) as total_file_creations,
        SUM(su_attempted) as total_root_attempts,
        SUM(num_root) as total_root_acceses
    FROM connections
    GROUP BY protocol_type, state
    ORDER BY 3 DESC
    """)

attack_stats.show()
```

Appending to spark\_analysis.py

```
[10]: %%writefile -a spark_analysis.py

ax[0].get_figure().savefig('report.png');
```

Appending to spark\_analysis.py

```
[11]: %%writefile -a spark_analysis.py

import google.cloud.storage as gcs
bucket = gcs.Client().get_bucket(BUCKET)
for blob in bucket.list_blobs(prefix='sparktodb/'):
    blob.delete()
bucket.blob('sparktodb/report.png').upload_from_filename('report.png')
```

Appending to spark\_analysis.py

```
[12]: %%writefile -a spark_analysis.py

connections_by_protocol.write.format("csv").mode("overwrite").save(
    "gs://{}/sparktodb/connections_by_protocol".format(BUCKET))
```

Appending to spark\_analysis.py

```
[13]: BUCKET_list = !gcloud info --format='value(config.project)'  
      BUCKET=BUCKET_list[0]  
      print('Writing to {}'.format(BUCKET))  
      !/opt/conda/anaconda/bin/python spark_analysis.py --bucket=$BUCKET
```

Writing to qwiklabs-gcp-02-6236be6b5a47

Setting default log level to "WARN".

To adjust logging level use `sc.setLogLevel(newLevel)`. For SparkR, use `setLogLevel(newLevel)`.

```
+-----+-----+  
|protocol_type| count|  
+-----+-----+  
|          icmp|283602|  
|           tcp|190065|  
|           udp| 20354|  
+-----+-----+  
  
+-----+-----+-----+-----+-----+-----+  
+-----+-----+-----+-----+-----+-----+  
+-----+  
|protocol_type|    state|total_freq|mean_src_bytes|mean_dst_bytes|mean_duration|  
total_failed_logins|total_compromised|total_file_creations|total_root_attempts|total_root_acceses|  
+-----+-----+-----+-----+-----+-----+  
+-----+-----+-----+-----+-----+-----+  
+-----+  
|          icmp|    attack|    282314|    932.14|    0.0|    0.0|  
0|              0|              0|    0.0|    0.0|  
|          tcp|    attack|    113252|    9880.38|    881.41|    23.19|  
57|              2269|              76|    1.0|  
152|  
|          tcp|no attack|    76813|    1439.31|    4263.97|    11.08|  
18|              2776|    459|    17.0|  
5456|  
|          udp|no attack|    19177|    98.01|    89.89|    1054.63|  
0|              0|    0|    0.0|    0.0|  
|          icmp|no attack|    1288|    91.47|    0.0|    0.0|  
0|              0|    0|    0.0|    0.0|  
|          udp|    attack|    1177|    27.5|    0.23|    0.0|  
0|              0|    0|    0.0|    0.0|  
+-----+-----+-----+-----+-----+-----+  
+-----+-----+-----+-----+-----+-----+  
+-----+  

```

Traceback (most recent call last):

File "spark\_analysis.py", line 70, in <module>

```
ax[0].get_figure().savefig('report.png');  
NameError: name 'ax' is not defined
```

```
[14]: !gsutil ls gs://$BUCKET/sparktodb/**
```

```
gs://qwiklabs-gcp-02-6236be6b5a47/sparktodb/spark_analysis.py
```

```
[15]: !gsutil cp spark_analysis.py gs://$BUCKET/sparktodb/spark_analysis.py
```

```
Copying file:///spark_analysis.py [Content-Type=text/x-python]...  
/ [1 files][ 2.7 KiB/ 2.7 KiB]  
Operation completed over 1 objects/2.7 KiB.
```

Copyright 2019 Google Inc. Licensed under the Apache License, Version 2.0 (the “License”); you may not use this file except in compliance with the License. You may obtain a copy of the License at <http://www.apache.org/licenses/LICENSE-2.0> Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an “AS IS” BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License.