

# Specify nested and repeated columns in table schemas

This page describes how to define a table schema with nested and repeated columns in BigQuery. For an overview of table schemas, see [Specifying a schema](/bigquery/docs/schemas) (/bigquery/docs/schemas).

## Define nested and repeated columns

To create a column with nested data, set the data type of the column to **RECORD** in the schema. A **RECORD** can be accessed as a **STRUCT** (/bigquery/docs/reference/standard-sql/data-types#struct\_type) type in standard SQL. A **STRUCT** is a container of ordered fields.

To create a column with repeated data, set the **mode** (/bigquery/docs/schemas#modes) of the column to **REPEATED** in the schema. A repeated field can be accessed as an **ARRAY** (/bigquery/docs/reference/standard-sql/data-types#array\_type) type in standard SQL.

A **RECORD** column can have **REPEATED** mode, which is represented as an array of **STRUCT** types. Also, a field within a record can be repeated, which is represented as a **STRUCT** that contains an **ARRAY**. An array cannot contain another array directly. For more information, see [Declaring an ARRAY type](/bigquery/docs/reference/standard-sql/data-types#declaring_an_array_type) (/bigquery/docs/reference/standard-sql/data-types#declaring\_an\_array\_type).

## Limitations

Nested and repeated schemas are subject to the following limitations:

**A schema cannot contain more than 15 levels of nested **RECORD** types.**

Columns of type `RECORD` can contain nested `RECORD` types, also called *child* records. The maximum nested depth limit is 15 levels. This limit is independent of whether the `RECORD`s are scalar or array-based (repeated).

## Example schema

The following example shows sample nested and repeated data. This table contains information about people. It consists of the following fields:

- `id`
- `first_name`
- `last_name`
- `dob` (date of birth)
- `addresses` (a nested and repeated field)
  - `addresses.status` (current or previous)
  - `addresses.address`
  - `addresses.city`
  - `addresses.state`
  - `addresses.zip`
  - `addresses.numberOfYearYears` (years at the address)

The JSON data file would look like the following. Notice that the addresses column contains an array of values (indicated by [ ]). The multiple addresses in the array are the repeated data. The multiple fields within each address are the nested data.

```
{ "id": "1", "first_name": "John", "last_name": "Doe", "dob": "1968-01-22", "addresses": [ { "status": "current", "address": "123 F:  
{ "id": "2", "first_name": "Jane", "last_name": "Doe", "dob": "1980-10-16", "addresses": [ { "status": "current", "address": "789 Ar
```

The schema for this table looks like the following:

```
[  
  {  
    "name": "id",  
    "type": "STRING",  
    "mode": "NULLABLE"  
  },  
  {  
    "name": "first_name",  
    "type": "STRING",  
    "mode": "NULLABLE"  
  },  
  {  
    "name": "last_name",  
    "type": "STRING",  
    "mode": "NULLABLE"  
  },  
  {  
    "name": "dob",  
    "type": "DATE",  
    "mode": "NULLABLE"
```

```
},  
{  
  "name": "addresses",  
  "type": "RECORD",  
  "mode": "REPEATED",  
  "fields": [  
    {  
      "name": "status",  
      "type": "STRING",  
      "mode": "NULLABLE"  
    },  
    {  
      "name": "address",  
      "type": "STRING",  
      "mode": "NULLABLE"  
    },  
    {  
      "name": "city",  
      "type": "STRING",  
      "mode": "NULLABLE"  
    },  
    {  
      "name": "state",  
      "type": "STRING",  
      "mode": "NULLABLE"  
    },  
    {  
      "name": "zip",  
      "type": "STRING",  
      "mode": "NULLABLE"  
    },  
  ],  
}
```

```

        "name": "numberOfYears",
        "type": "STRING",
        "mode": "NULLABLE"
    }
]
}
]

```

## Specifying the nested and repeated columns in the example

Console (#console) bq (#bq) Go (#go) Java (#java) Node.js (#node.js) Python (#python)

Before trying this sample, follow the Python setup instructions in the [BigQuery quickstart using client libraries](#) (/bigquery/docs/quickstarts/quickstart-client-libraries). For more information, see the [BigQuery Python API reference documentation](#) (https://googleapis.dev/python/bigquery/latest/index.html).

[View on GitHub](https://github.com/googleapis/python-bigquery/blob/35627d145a41d57768f19d4392ef235928e00f72/docs/snippets.py) (https://github.com/googleapis/python-bigquery/blob/35627d145a41d57768f19d4392ef235928e00f72/docs/snippets.py)

```

# from google.cloud import bigquery
# client = bigquery.Client()
# project = client.project
# dataset_ref = bigquery.DatasetReference(project, 'my_dataset')

schema = [
    bigquery.SchemaField("id", "STRING", mode="NULLABLE"),
    bigquery.SchemaField("first_name", "STRING", mode="NULLABLE"),
    bigquery.SchemaField("last_name", "STRING", mode="NULLABLE"),

```

```
bigquery.SchemaField("dob", "DATE", mode="NULLABLE"),
bigquery.SchemaField(
    "addresses",
    "RECORD",
    mode="REPEATED",
    fields=[
        bigquery.SchemaField("status", "STRING", mode="NULLABLE"),
        bigquery.SchemaField("address", "STRING", mode="NULLABLE"),
        bigquery.SchemaField("city", "STRING", mode="NULLABLE"),
        bigquery.SchemaField("state", "STRING", mode="NULLABLE"),
        bigquery.SchemaField("zip", "STRING", mode="NULLABLE"),
        bigquery.SchemaField("numberOfYears", "STRING", mode="NULLABLE"),
    ],
),
]
table_ref = dataset_ref.table("my_table")
table = bigquery.Table(table_ref, schema=schema)
table = client.create_table(table) # API request

print("Created table {}".format(table.full_table_id))
```

## Modify nested and repeated columns

After you add a nested column or a nested and repeated column to a table's schema definition, you can modify the column as you would any other type of column. BigQuery natively supports several schema changes such as adding a new nested field to a record or relaxing a nested field's mode. For more information, see [Modifying table schemas](https://cloud.google.com/bigquery/docs/managing-table-schemas) (/bigquery/docs/managing-table-schemas).

Additionally, you can manually modify a schema definition that includes nested and repeated columns. For more information, see [Manually changing table schemas](/bigquery/docs/manually-changing-schemas) (/bigquery/docs/manually-changing-schemas).

## When to use nested and repeated columns

BigQuery performs best when your data is denormalized. Rather than preserving a relational schema such as a star or snowflake schema, denormalize your data and take advantage of nested and repeated columns. Nested and repeated columns can maintain relationships without the performance impact of preserving a relational (normalized) schema.

For example, a relational database used to track library books would likely keep all author information in a separate table. A key such as `author_id` would be used to link the book to the authors.

In BigQuery, you can preserve the relationship between book and author without creating a separate author table. Instead, you create an author column, and you nest fields within it such as the author's first name, last name, date of birth, and so on. If a book has multiple authors, you can make the nested author column repeated.

BigQuery supports loading nested and repeated data from source formats that support object-based schemas, such as JSON files, Avro files, Firestore export files, and Datastore export files.

## Table security

To control access to tables in BigQuery, see [Introduction to table access controls](/bigquery/docs/table-access-controls-intro) (/bigquery/docs/table-access-controls-intro).

## Next steps

- To insert and update rows with nested and repeated columns, see [Data manipulation language syntax](https://cloud.google.com/bigquery/docs/reference/standard-sql/dml-syntax) (/bigquery/docs/reference/standard-sql/dml-syntax).

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 4.0 License](https://creativecommons.org/licenses/by/4.0/) (https://creativecommons.org/licenses/by/4.0/), and code samples are licensed under the [Apache 2.0 License](https://www.apache.org/licenses/LICENSE-2.0) (https://www.apache.org/licenses/LICENSE-2.0). For details, see the [Google Developers Site Policies](https://developers.google.com/site-policies) (https://developers.google.com/site-policies). Java is a registered trademark of Oracle and/or its affiliates.

Last updated 2021-08-17 UTC.